

#Import necessary libraries

```
In [2]: ▶ import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import GridSearchCV
```

#Load the dataset

```
In [3]: ▶ data = pd.read_csv("C:\\Users\\HP\\Desktop\\car_data.csv")
data
```

Out[3]:

	User ID	Age	AnnualSalary	Purchased
0	385	35	20000	0
1	681	40	43500	0
2	353	49	74000	0
3	895	40	107500	1
4	661	25	79000	0
...
995	863	38	59000	0
996	800	47	23500	0
997	407	28	138500	1
998	299	48	134000	1
999	687	44	73500	0

#Perform data preprocessing

```
In [4]: x = data.drop('Purchased', axis=1)
x
```

Out[4]:

	User ID	Age	AnnualSalary
0	385	35	20000
1	681	40	43500
2	353	49	74000
3	895	40	107500
4	661	25	79000
...
995	863	38	59000
996	800	47	23500
997	407	28	138500
998	299	48	134000
999	687	44	73500

1000 rows × 3 columns

```
In [5]: y=data['Purchased']
y
```

Out[5]:

0	0
1	0
2	0
3	1
4	0
...	..
995	0
996	0
997	1
998	1
999	0

Name: Purchased, Length: 1000, dtype: int64

```
In [6]: ▶ #Check for the missing data  
missing_data = data.isnull()  
  
#Count missing values in each column  
missing_count = missing_data.sum()  
  
#Display the count of missing values  
print("Missing Data")  
print(missing_count)
```

Missing Data

User ID 0

Age 0

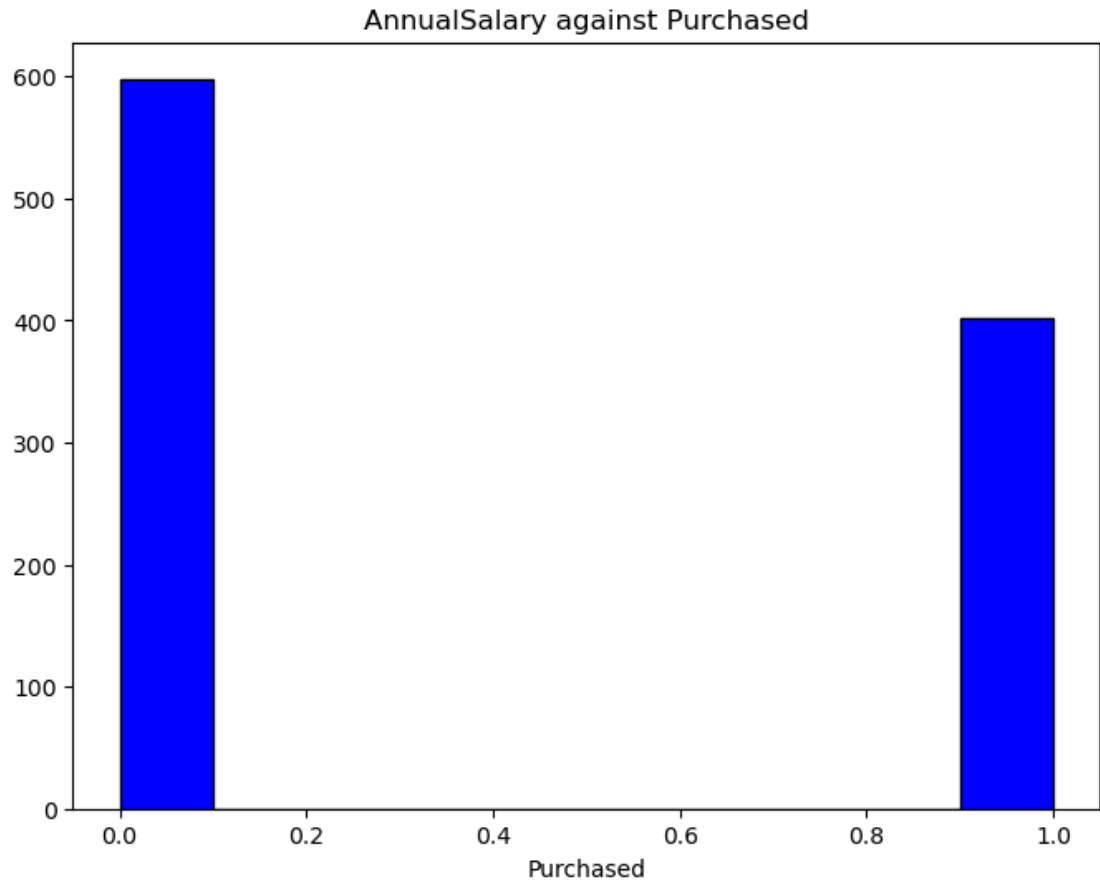
AnnualSalary 0

Purchased 0

dtype: int64

In [7]: `#Visualising the graphics`

```
import matplotlib.pyplot as plt
plt.figure(figsize=(8,6))
plt.hist(y, bins=10, color='blue', edgecolor='black')
plt.ylabel('')
plt.xlabel('AnnualSalary')
plt.xlabel('Purchased')
plt.title('AnnualSalary against Purchased')
plt.show()
```



In [8]: `#Splitting the data`
`#Import the library`
`from sklearn.model_selection import train_test_split`

In [9]: `#Split the data into training and testing sets`
`x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_st`

```
In [10]: #Standardising the data
#Importing library
from sklearn.preprocessing import StandardScaler

#Initialise the StandardScaler
scaler = StandardScaler()

#Fit the scaler to the training data and transform the training data
x_train_scaled = scaler.fit_transform(x_train)

#Transform the testing data using the same scaler
x_test_scaled = scaler.transform(x_test)
```

```
In [11]: #Building the model
```

```
In [12]: from sklearn.linear_model import LogisticRegression

#Initialise the Logistic Regression model
logistic_regression_model = LogisticRegression()
```

```
In [13]: #Fitting the model on the training data
logistic_regression_model.fit(x_train_scaled,y_train)
```

```
Out[13]: LogisticRegression
LogisticRegression()
```

```
In [14]: #Preddicting the target variable for the testing data
y_pred = logistic_regression_model.predict(x_test_scaled)
y_pred
```

```
Out[14]: array([0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1,
                0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0,
                1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
                0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,
                0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
                0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0,
                0, 1], dtype=int64)
```

```
In [15]: ▶ #Getting coefficient
coefficients = logistic_regression_model.coef_

#Getting intercept
intercept = logistic_regression_model.intercept_

#Displaying coefficient and intercept
print("Coefficients:", coefficients)
print("intercept:", intercept)

Coefficients: [[-0.01217401  2.2094714  1.17459755]]
intercept: [-0.82747317]
```

```
In [16]: ▶ #Evaluating the model
```

```
In [17]: ▶ from sklearn.metrics import accuracy_score, precision_score, recall_score,
```

```
In [18]: ▶ #Checking for uniques values in y_test and y_pred
unique_y_test = np.unique(y_test)
unique_y_pred = np.unique(y_pred)

print("Unique values in y_test:", unique_y_test)
print("Unique values in y_pred:", unique_y_pred)
```

```
Unique values in y_test: [0 1]
Unique values in y_pred: [0 1]
```

```
In [19]: ▶ # compute accuracy
accuracy = accuracy_score(y_test, y_pred)
accuracy
```

```
Out[19]: 0.815
```

```
In [20]: ▶ #Compute precision
precision = precision_score(y_test, y_pred)
precision
```

```
Out[20]: 0.8805970149253731
```

```
In [21]: ▶ #Compute recall
recall = recall_score(y_test, y_pred)
recall
```

```
Out[21]: 0.6704545454545454
```

```
In [22]: ▶ #Convert class labels to binary format
y_test_binary = np.where(y_test ,1,0)
y_pred_binary = np.where(y_pred ,1,0)

#Compute f1-score
f1 = f1_score(y_test_binary, y_pred_binary)

#Print f1-score
print("f1-score:", f1)
```

f1-score: 0.7612903225806451

```
In [23]: ▶ #Compute predicted probabilities
y_prob = logistic_regression_model.predict_proba(x_test_scaled)

#Extract the probabilities for the positive class
y_prob_positive = y_prob[:,1]

#Compute ROC-AUC score
roc_auc = roc_auc_score(y_test_binary,y_prob_positive)

#Print ROC-AUC score
print("ROC-AUC score:", roc_auc)
```

ROC-AUC score: 0.8933644480519481

```
In [42]: ▶ #Display the evaluation metrics
print("Accuracy:", accuracy)
print("Precision:",precision)
print("Recall:",recall)
print("F1-score:",f1)
print("ROC-AUC score:",roc_auc)
```

Accuracy: 0.815
Precision: 0.8805970149253731
Recall: 0.6704545454545454
F1-score: 0.7612903225806451
ROC-AUC score: 0.8933644480519481

MODEL OPTIMISATION

```
In [53]: ▶ from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
```

```
In [54]: ▶ model = LogisticRegression()
model
```

```
Out[54]: ▾ LogisticRegression
LogisticRegression()
```

```
In [55]: ▶ #Define the Grid of hyperparameters
param_grid = {'penalty': ['l1', 'l2'],
              'C': [0.1, 1],
              'solver': ['liblinear', 'saga']}
}
```

```
In [56]: ▶ #Initailise GridSearchCV
grid_search = GridSearchCV(model, param_grid, cv=5, scoring='accuracy')
```

```
In [57]: ▶ #Perform grid search to find the best hyperparameters
grid_search.fit(x_train_scaled, y_train)
```

```
Out[57]: ▸ GridSearchCV
▸ estimator: LogisticRegression
  ▸ LogisticRegression
```

```
In [58]: ▶ #Getting the best hyperparameters
best_params = grid_search.best_params_
best_params
```

```
Out[58]: {'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}
```

```
In [59]: ▶ #Initialise logistic regression model with the best hyperparameters
best_logistic_regression_model = LogisticRegression(**best_params)
```

```
In [60]: ▶ #Fit the model on the training data
best_logistic_regression_model.fit(x_train_scaled, y_train)
```

```
Out[60]: ▾ LogisticRegression
LogisticRegression(C=0.1, penalty='l1', solver='liblinear')
```

#Split the dataset into training and testing sets


```
In [61]: #Predict the target variables for testing data
y_pred_best = best_logistic_regression_model.predict(x_test_scaled)
y_pred_best
```

```
Out[61]: array([0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1,
0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0,
1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,
0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0,
0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0,
0, 1], dtype=int64)
```

```
In [62]: #Evaluate the model
accuracy_best = accuracy_score(y_test, y_pred_best)
precision_best = precision_score(y_test, y_pred_best)
recall_best = recall_score(y_test_binary, y_pred_binary)
f1_best = f1_score(y_test_binary, y_pred_binary)
y_prob_best = best_logistic_regression_model.predict_proba(x_test_scaled)
y_prob_postive_best = y_prob_best[:, 1]
roc_auc_best = roc_auc_score(y_test_binary, y_prob_postive_best)

#Display the evaluation metrics for optimised model
print("Optimized Model Evaluation Metrics:")
print("Accuracy:", accuracy_best)
print("Precision:", precision_best)
print("Recall:", recall_best)
print("F1-score:", f1_best)
print("ROC-AUC Score:", roc_auc_best)
```

```
Optimized Model Evaluation Metrics:
Accuracy: 0.825
Precision: 0.8840579710144928
Recall: 0.6704545454545454
F1-score: 0.7612903225806451
ROC-AUC Score: 0.8947849025974026
```